

Lecture 3: PAC Learning

Instructor: Santosh Vempala

Lecture date: 9/06,9/11

1 Learning the Class of Conjunctions

Consider data of the form $\{x_1, \dots, x_n\} \in \{0, 1\}^n$. A *conjunction* is a function on some subset of the variables $\{x_1, x_2, \dots, x_n, \bar{x}_1, \bar{x}_2, \dots, \bar{x}_n\}$, where $\bar{x}_i = 1 - x_i$. It maps x to 1 if every variable in the subset is 1, and 0 otherwise. For example:

$$\begin{aligned} h(x) &= x_1 \wedge x_2 \wedge \bar{x}_5 \\ h(x) &= \bar{x}_3 \wedge \bar{x}_4 \\ h(x) &= x_1 \wedge x_2 \wedge \dots \wedge x_r \end{aligned}$$

The size of the hypothesis class is 3^n ; for each i , either x_i is in the conjunction, \bar{x}_i is in the conjunction, or both are absent.

Any such conjunction can be represented as a linear halfspace, so the tools for learning halfspaces seen in previous lectures apply without modification. Simply sum the variables in the conjunction, and set the threshold to be the number of variables. Using the examples above:

$$\begin{aligned} h(x) &= (x_1 + x_2 + 1 - x_5 \geq 3) \\ h(x) &= (1 - x_3 + 1 - x_4 \geq 2) \\ h(x) &= (x_1 + x_2 + \dots + x_r \geq r) \end{aligned}$$

If $x_i \in \{0, 1\}$, these two ways of writing the function are equivalent.

1.1 Learn-Conj

Let \mathcal{D} be a distribution over $\{0, 1\}^n$. Assume that each example is chosen independently from this distribution. Consider the following algorithm for learning an unknown conjunction.

Algorithm 1 LEARNCONJ

```

Maintain two sets of indices  $S = \{x_1, x_2, \dots, x_n\}, \bar{S} = \{\bar{x}_1, \bar{x}_2, \bar{x}_n\}$ 
for each input  $x^{(i)}$  do
  if  $l(x^{(i)}) = 1$  then
    For each  $j \in [n]$ , remove  $\bar{x}_j$  from  $\bar{S}$  if  $x_j^{(i)} = 1$ , and remove  $x_j$  from  $S$  if  $x_j^{(i)} = 0$ 
  else
    Do nothing
  end if
end for
Output the conjunction of the remaining literals in  $S$  and  $\bar{S}$ .

```

At any point, the hypothesis $h(x) = (\bigwedge_{i \in S} x_i) \wedge (\bigwedge_{i \in \bar{S}} \bar{x}_i)$ is consistent with all examples seen so far.

Theorem 1. With $m \geq \frac{1}{\epsilon}(n \log 3 + \log \frac{1}{\delta})$ examples, with probability $1 - \delta$, LEARNCONJ will be correct on a new example from \mathcal{D} with probability $1 - \epsilon$.

Proof. Let $h \in H$ be a hypothesis with $\Pr_{x \sim \mathcal{D}}(h(x) \neq l(x)) > \epsilon$. The probability that h is not eliminated after m examples is at most $(1 - \epsilon)^m$.

The probability that any such ‘bad h ’ survives after m examples is at most $3^n(1 - \epsilon)^m$.

Setting $3^n(1 - \epsilon)^m < \delta$ and taking the log of both sides, we see that $m \geq \frac{1}{\epsilon}(n \log 3 + \log \frac{1}{\delta})$ suffices. □

1.2 Learning Decision Lists

A *decision list* is a generalization of the set of conjunctions. It consists of a series of if-else statements, returning true/false at each step. E.g.,

If $x_1 = 1$, return False; else if $x_4 = 0$, return True; else, return False.

Any conjunction can be written as a decision list. E.g., $x_1 \wedge x_2 \wedge \dots \wedge x_r$ is equivalent to:

If $x_1 = 0$, return False; else if $x_2 = 0$, return False; ... else if $x_r = 0$, return False; else, return True

The number of possible decision lists is at most $4^n n!$ (choosing 0/1 for the check and return statement, and every possible ordering of the indices).

Suppose we have an algorithm LEARNDL that maintains a decision list h which is consistent with all examples seen so far. Then, the following theorem applies:

Theorem 2. *With $m \geq \frac{c}{\epsilon}(n \log n + \log \frac{1}{\delta})$ examples, with probability $1 - \delta$, LEARNDL will be correct on a new example from \mathcal{D} with probability $1 - \epsilon$.*

Proof. The argument from Theorem 1 does not use any special properties of conjunctions or the algorithm. We can apply the same logic here.

Let $h \in H$ be a hypothesis with $Pr_{x \sim \mathcal{D}}(h(x) \neq l(x)) > \epsilon$. The probability that h is not eliminated after m examples is at most $(1 - \epsilon)^m$.

The probability that any such ‘bad h ’ survives after m examples is at most $4^n n!(1 - \epsilon)^m$. (This time, substituting the size of the class of decision lists!)

Setting $4^n n!(1 - \epsilon)^m < \delta$ and taking the log of both sides, we see that $m \geq O\left(\frac{1}{\epsilon}\right)(n \log n + \log \frac{1}{\delta})$ suffices. \square

In general, for an algorithm that learns a consistent hypothesis from a class H , $m \geq \frac{c}{\epsilon}(\log |H| + \log \frac{1}{\delta})$ examples are sufficient to achieve this guarantee.

2 (ϵ, δ) -PAC-Learning

Definition 3. *Given iid samples from an unknown distribution \mathcal{D} , labeled by a hypothesis ℓ from a hypothesis class H , an algorithm \mathcal{A} (ϵ, δ) -PAC-learns the hypothesis class H if, with probability $1 - \delta$, it produces a hypothesis $h \in H$ that correctly classifies a random sample from \mathcal{D} with probability at least $1 - \epsilon$, i.e.,*

$$Pr_{x \in \mathcal{D}}(h(x) = \ell(x)) \geq 1 - \epsilon.$$

PAC stands for Probably Approximately Correct. Intuitively, it states that the algorithm \mathcal{A} is *probably* successful (producing a good h with probability $1 - \delta$). A good h is *approximately correct* (classifying samples from \mathcal{D} correctly with probability $1 - \epsilon$).

Note that this definition says nothing about the efficiency of the algorithm. Usually the efficiency is judged by time or sample complexity (i.e., the time it takes to process a sample, and the number of samples that it needs).

2.1 PAC-learning Halfspaces

In section 1, we showed a heuristic for learning an arbitrary hypothesis class H . However, the number of samples needed depends on $\log |H|$. For the case of halfspaces, H is infinite, so we must use another method.

Suppose, $\|w^*\|_2 \leq 1$, $\max_x \|x\|_2 \leq 1$, and $\max_x |\langle w^*, x \rangle| \geq \gamma$. At any time, define the set of consistent hypothesis.

$$W = \{w : \|w\|_2 \leq 1, \forall i \langle w, l(x^{(i)})x^{(i)} \rangle \geq \gamma\}$$

For the next sample $x^{(i+1)}$, divide W into two cases:

$$W_+ = \{w \in W : \langle w, x^{(i+1)} \rangle \geq 0\}$$

$$W_- = \{w \in W : \langle w, x^{(i+1)} \rangle < 0\}$$

By predicting the set with the larger volume, we guarantee that either we are correct, or the size of W is reduced by at least half.

Note that because the margin is γ , we only need to reduce W to within a ‘cap’ of angle γ . The set is

$$V_\gamma = \{w : \|w\|_2 = 1, \langle w, w^* \rangle \geq 1 - \gamma^2\}$$




Figure 1: The set of points on the sphere where $\langle w, w^* \rangle \geq 1 - \gamma^2$

This set is illustrated in red in Figure 1. All vectors in V_γ will correctly classify examples that obey the margin.

Theorem 4. *The number of steps required to find a consistent classifier is $O(n \log \gamma)$.*

Proof. The volume of W is reduced by at least half at each step; the initial set is $S_{n-1} = \{w : \|w\| = 1\}$. So, we need to bound $\log_2 \frac{\text{Vol}(S_{n-1})}{\text{Vol}(V_\gamma)}$.

To do this, we can write both sets as an integral over the first coordinate. Note that when we fix $x_1 = t$, then by the equation for the sphere:

$$\sum_{i=1}^n x_i^2 = 1 \rightarrow \sum_{i=2}^n x_i^2 = 1 - t^2$$

Hence, the cap formed by intersecting the sphere with the plane $x_1 = t$ is a sphere of radius $\sqrt{1 - t^2}$

$$\text{vol}(S_{n-1}) = 2 \int_0^1 (1 - t^2)^{(n-2)/2} \text{Vol}(S_{n-2}) dt$$

$$\text{vol}(V_\gamma) = \int_{\sqrt{1-\gamma^2}}^1 (1 - t^2)^{(n-2)/2} \text{Vol}(S_{n-2}) dt$$

This integral is not straight-forward to evaluate; fortunately, we only need an asymptotic bound, so we can make some simplifications.

Simplifying this integral

$$\begin{aligned} \int_{\sqrt{1-\gamma^2}}^1 (1-t^2)^{(n-2)/2} dt &\geq \int_{\sqrt{1-\gamma^2}}^1 (\gamma^2)^{(n-2)/2} dt \\ &= (1 - \sqrt{1-\gamma^2}) \gamma^{n-2} \\ &= c^n \gamma^n \text{ for a constant } c \end{aligned}$$

Since $\int_0^1 (1-t^2)^{(n-2)/2} dt \leq 1$, the ratio is

$$\frac{\text{Vol}(S_{n-2}) \int_{\cos(\gamma)}^1 (1-t^2)^{(n-2)/2} dt}{2\text{Vol}(S_{n-2}) \int_0^1 (1-t^2)^{(n-2)/2} dt} \geq c^n \gamma^n$$

Taking the logarithm gives us the theorem statement. \square

3 Ways to label points

Definition 5. For a concept class \mathcal{H} and an integer $m > 0$, let $\mathcal{H}(m)$ be the maximum number of distinct ways a set of m points can be labelled using concepts in \mathcal{H}

Example: Intervals in 1-D

Let $\mathcal{H} = \{[a, b] : a, b \in \mathbb{R}, a < b\}$; the set of all intervals on the real line.

As illustrated in Figure 2, $\mathcal{H}(2) = 4$, since the two points shown in the plot can be labelled in all four ways. However, $\mathcal{H}(3) = 7$, since there is a sequence of labels $(+ - +)$ which cannot be achieved with a single interval (but all other labellings are possible!)




Figure 2: The four ways to label a set of points on the real-line with a single interval




Figure 3: A set of three points with an assignment of labels that can't be achieved by a single interval

Rectangles in 2D

Let $\mathcal{H} = \{[a, b] \times [c, d] : a, b, c, d \in \mathbb{R}, a < b, c < d\}$; the set of all rectangles in \mathbb{R}^2 .

We have $\mathcal{H}(m) \leq (m+1)^4$. There are $(m+1)^2$ ways to pick the vertical lines, and $(m+1)^2$ ways to pick the horizontal lines.

Halfspaces in \mathbb{R}^n

Note that in \mathbb{R}^2 , defining a separating line requires 2 points. This generalizes to \mathbb{R}^n , where it takes n points to define a separating hyperplane. We can choose our labelling by picking n points out of the sample to define the plane.

This gives:

$$\mathcal{H}(m) \leq \binom{m}{n} * 2 \leq m^n \quad (1)$$

Theorem 6. *Suppose we have a true label $h \in \mathcal{H}$. For (ϵ, δ) -PAC-Learning, we can output an $\tilde{h} \in \mathcal{H}$ that is consistent in expectation provided that:*

$$m \geq \frac{c}{\epsilon} \left(\log \mathcal{H}(2m) + \log \frac{1}{\delta} \right)$$

Corollary 7. *The sample complexity of (ϵ, δ) -PAC-Learning halfspaces in \mathbb{R}^n is $\frac{c}{\epsilon} \left(n \log \frac{1}{\epsilon} + \log \frac{1}{\delta} \right)$.*

Proof. To argue this, we can substitute the bound given in Equation 1 into the theorem statement.

$$m \geq \frac{c}{\epsilon} \left(\log(2m)^n + \log \frac{1}{\delta} \right)$$

Solving for m leads to the corollary. □